Chapter 1

Variant's generation

RQ1. To what extent it is possible to artificially create variants for WebAssembly programs?

In this chapter, we investigate whether we can artifically create program variants through semantically equivalent code transformations. We propose a framework to create program variants that are functionally equivalent to their original, yet exhibit a different behavior. We propose to retarget a superoptimizer, using its exhaustive searching strategy for providing semantically equivalent code transformations. Furthermore, we present a framework, able to generate program variants that can be successfully compiled to WebAssembly. The presented methodology and transformation tool, CROW, are contributions to this thesis. We evaluate the usage of CROW on a corpus of open-source and nature diverse programs. We sum up the key insights taken from this evaluation.

1.1 CROW

In this section we describe CROW, a tool tailored to create semantically equivalent variants out of a single program, either C/C++ code or LLVM bitcode. CROW is part of the contributions of this thesis. In Figure 1.1, we describe the workflow of CROW to create program variants.

CROW synthesizes program variants to be WebAssembly binaries. We assume that the programs are generated through the LLVM compilation pipeline. This assumption is supported by the work of Lehman et al. [] and the fact that Wasm programs are generated in the 60% of the times by the LLVM toolchain.

During the *exploration* stage, CROW takes as input a C/C++ programs or LLVM bitcodes and produces a set of unique, diversified LLVM bitcodes that can be later ported to WebAssembly. Figure 1.1 shows the stages of this workflow. The workflow starts with compiling the input program into LLVM bitcode using clang
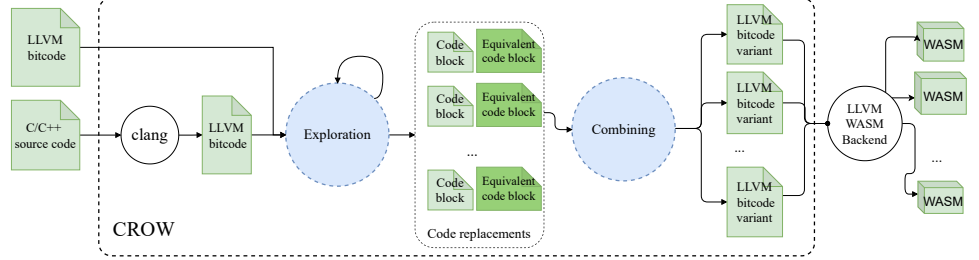
Figure 1.1: CROW workflow to generate program variants. CROW takes C/C++ source codes or LLVM bitcodes to look for code blocks that can be replaced by semantically equivalent code and generates program variants by combining them.

if it comes from source code. Then, CROW analyzes the bitcode to synthesize a set of candidate code replacements. In the following we enunciate the definitions we use along this work for code block, functional equivalence and code replacement.

Definition 1 *Block (based on Aho et al. [?]): Let P be a program. A block B is a grouping of declarations and statements in P inside a function F.*

Definition 2 *Functional equivalence modulo program state (based on Le et al. [?]): Let B_1 and B_2 be two code blocks according to Definition 1. We consider the program state before the execution of the block, S_i, as the input and the program state after the execution of the block, S_o, as the output. B_1 and B_2 are functionally equivalent if given the same input S_i both codes produce the same output S_o.*

Definition 3 *Code replacement: Let P be a program and T a pair of code blocks (B_1, B_2). T is a candidate code replacement if B_1 and B_2 are both functionally equivalent as defined in Definition 2. Applying T to P means replacing B_1 by B_2. The application of T to P produces a program variant P' which consequently is functionally equivalent to P.*

We address the *exploration* stage by retargeting a superoptimizer for LLVM, using its subset of the LLVM intermediate representation. CROW operates at function level, taking the functions inside the LLVM bitcode module as individual instances to analyze and diversify. The retargeted superoptimizer is in charge of finding the potential places in the original code functions where a replacement can be applied. Also, it makes the formal verification of Definition 2 using a theorem prover. We prevent the superoptimizer from synthesizing instructions that have no correspondence in WebAssembly for sake of reducing the searching space for equivalent program variants. Besides, we disable all optimizations in the Web-Assembly LLVM backend that could reverse the superoptimizer transformations, such as constant folding and instructions normalization.

Finally, in the *combining* stage, CROW combines the candidate code replacements to generate different LLVM bitcode variants, selecting and combining the code replacements. CROW applies the candidate code replacements to the original program to produce a LLVM bitcode variants. Then, this bitcode is compiled into a WebAssembly binary if requested. CROW generates the variants from all possible combinations of code replacements as the power set over all code replacements.

1.1.1 Example

Let us illustrate how CROW works with the simple example code in Listing 1.1. The f function calculates the value of $2 * x + x$ where x is the input for the function. This code is compiled by CROW generating the intermediate LLVM bitcode in the left most part of Listing 1.2. As the right-most part of Listing 1.2 shows, CROW found two code blocks to look for variants.

Listing 1.1: C function that calculates the quantity $2x + x$

```
int f(int x) {
    return 2 * x + x;
}
```

Listing 1.2: LLVM's intermediate representation program, its extracted code blocks and replacement candidates. Gray highlighted lines represent original code block, green for candidate code replacements.

```
define i32 @f(i32) {          Replacement candidates for    Replacement candidates for
                                     code_block_1                   code_block_2
       code block 2
       code block 1           %2 = mul nsw i32 %0,2          %3 = add nsw i32 %0,%2
   %2 = mul nsw i32 %0,2
   %3 = add nsw i32 %0,%2      %2 = add nsw i32 %0,%0         %3 = mul nsw %0, 3:i32

   ret i32 %3                 %2 = shl nsw i32 %0, 1:i32
 }

 define i32 @main() {
 %1 = tail call i32 @f(i32
       10)
 ret i32 %1
 }
```

CROW, in the exploration stage detects 2 code blocks, **code_block_1** and **code_block_2** as the enclosing boxes in the left most part of Listing 1.2 show. CROW synthesizes $2 + 1$ candidate code replacements for each code block respectively as the green highlighted lines show in the right most parts of Listing 1.2. The baseline strategy of CROW to generate variants is to generate variants out of all

Listing 1.3: Candidate code replacements combination. Orange highlighted code illustrate replacement candidate overlapping.

```
%2 = mul nsw i32 %0,2            %2 = mul nsw i32 %0,2
%3 = add nsw i32 %0,%2           %3 = mul nsw %0, 3:i32

%2 = add nsw i32 %0,%0           %2 = add nsw i32 %0,%0
%3 = add nsw i32 %0,%2           %3 = mul nsw %0, 3:i32

%2 = shl nsw i32 %0, 1:i32       %2 = shl nsw i32 %0, 1:i32
%3 = add nsw i32 %0,%2           %3 = mul nsw %0, 3:i32
```

possible combinations of the candidate code replacements, *i.e.,* uses the power set of all candidate code replacements.

In the example, the power set is the cartesian product of the found candidate code replacements for each code block including the original code blocks, as Listing 1.3 shows. The power set size results in 6 potential function variants. Yet, the generation stage would eventually generate 4 variants from the original WebAssembly binary. This gap between the number of potential and the actual number of variants is a consequence of the redundancy among the bitcode variants when composing into one. In other words, if the replaced code does not involve other code blocks with replacements, all possible combinations having it will be in the end the same program. In the example case, replacing `code_block_2` by `mul nsw %0, 3`, turns `code_block_1` into dead code, thus, later replacements generate the same program variants. The right most part of Listing 1.3 illustrates how for three different combinations, CROW produces the same variant. We call this phenomenon a candidate code replacement overlapping.

One might think that a proper heuristic could be implemented to avoid such case of overlapping. We have found easier and faster to generate the variants with the replacements combination and checking their uniqueness after the program variant is compiled. This prevents us of having an expensive checking for overlapping inside the CROW code. Still, this phenomenon calls for later optimizations in future works.

The final result can be appreciated in Listing 1.4. CROW generated 4 statically different Wasm files.

1.2   Evaluation

We use CROW, the tool described at section 1.1 to answer RQ1. In this section we describe the corpora of original programs that we pass to CROW for sake of variants generation. Besides, we describe our metrics, and we finalize by discussing the results.

Listing 1.4: Wasm program variants generated from program Listing 1.1.

```
func $f (param i32) (result i32)        func $f (param i32) (result i32)
  local.get 0                             local.get 0
  i32.const 2                             i32.const 1
  i32.mul                                 i32.shl
  local.get 0                             local.get 0
  i32.add                                 i32.add


func $f (param i32) (result i32)        func $f (param i32) (result i32)
  local.get 0                             local.get 0
  local.get 0                             i32.const 3
  i32.add                                 i32.mul
  local.get 0
  i32.add
```

1.2.1 Corpora

We answer RQ1 with two corpora of programs appropriate for our experiments. The first corpus, **CROW prime**, is part of the CROW contribution []. The second corpus, **MEWE prime**, is part of the MEWE contribution []. In Table 1.1 we summarize the selection criteria, and we mention each corpus properties. With both corpora we evaluate CROW with a total of $303+7$ programs, containing $303+1902$ functions. To assess the ability of our approach to generate WebAssembly binaries that are statically different, we compute the number of unique variants generated by CROW for each original program. We compare the WebAssembly program and its variant using the `md5` hash of each program bytestream as a metric for uniqueness.

1.2.2 Setup and evaluation

CROW's workflow synthesizes program variants with an enumerative strategy. This means that all possible programs that can be generated for a given language (LLVM in the case) are constructed and verified for equivalence. There are two parameters to control the size of the search space and hence the time required to traverse it. On one hand, one can limit the size of the variants. On the other hand, one can limit the set of instructions that are used for the synthesis. In our experiments, we use between 1 instruction (only additions) and 60 instructions (all supported instructions in the synthesizer).

These two configuration parameters allow the user to find a trade-off between the amount of variants that are synthesized and the time taken to produce them. In Table 1.2 we listed the configuration for both corpora. For the current evaluation, given the size of the corpus, we set the exploration time to 1 hour maximum per function for **CROW PRIME**, for a total of 303 hours CROW executions. In the case of **MEWE prime** we set the timeout to 5 minutes per function in the exploration stage. We set all 60 supported instructions in CROW for both **CROW prime** and **MEWE primer** corpora.

Corpus name	Selection criteria	Corpus Description
CROW prime	To build the corpus used in CROW, we take programs from the Rosetta Code project[1]. We first collect all C programs from Rosetta Code, which represents 989 programs as of 01/26/2020. We then apply a number of filters: the programs should successfully compile, they should not require user inputs, the programs should terminate and should not provide in non-deterministic results. The result of the filtering is a corpus of 303 C programs	All programs have a single function in terms of source code. These programs range from 7 to 150 lines of code and solve a variety of problems, from the *Babbage* problem to *Convex Hull* calculation.
MEWE prime	We select two mature and typical edge-cloud computing projects for this corpus. The projects are selected based on their suitability for diversity synthesis with CROW, *i.e.,* the projects should have the ability to collect their modules in LLVM intermediate representation The selected projects are: **libsodium**, an encryption, decryption, signature and password hashing library which can be ported to WebAssembly and **qrcode-rust**, a QrCode and MicroQrCode generator written in Rust. We then filter out 5 and 2 endpoints[2] respectively, for which we select their involved functions.	The evaluated projects contain in total 1902 functions, 62 for libdosium and 1840 for qrcode-rust. The function range between 10 ad 127700 lines of code.

Table 1.1: Corpora description. The table is composed by the name of the corpus, the selection criteria and the stats the programs in each corpus.

CORPUS	Exploration timeout	Max. instructions
CROW prime	1h	60
MEWE prime	5m	60

Table 1.2: CROW tweaking for variants generation. The table is composed by the name of the corpus, the timeout parameter and the maximum number of instructions allowed in the synthesis process.

1.3   Results

We summarize the results in Table 1.3 CROW produces at least one unique program variant for 239/303 programs for **CROW prime** with 1h for timeout. For the rest of the programs (64/303), the timeout is reached before CROW can find any valid variant. In the case of **MEWE prime**, CROW produces equivalent variants for

48/1902 original programs with 5 minutes per function as timeout. The rest of the programs resulted timeout before finding function variants or produced none.

CORPUS	#Functions	# Diversified	# NonDiversified	# Variants
CROW prime	303	**239**	64	1906
MEWE prime	1902	48	1854	**4670**

Table 1.3: General diversification results. The table is composed by the name of the corpus, the number of functions, the number of succesfully diversified functions, the number of non-diversified functions and the number of unique variants.

1.3.1 Challenges for automatic diversification

CROW generates variants for functions in both corpora, however, we have observed a remarkable difference between the number of successfully diversified functions versus the number of failed-to-diversify functions, as it can be appreciated in Table 1.3. CROW succesfully diversified approx. 79% and 2.5% of the original functions for **CROW prime** and **MEWE prime** respectively. On the other hand, CROW generated more variants for **MEWE prime**, 4670 program variants for 48 original programs. Not surprisingly, to set the timeout affects the capacity of CROW for diversification. On the other hand, a low timeout for exploration gives CROW more power of combining code replacements. This can be appreciated in the last column of the table, where for a lower number of diversified functions CROW created, overall, more variants.

Moreover, we look at the cases that yield a few variants per function. There is no direct correlation between the number of identified code for replacement and the number of unique variants. We manually analyze programs that include a significant number of potential places for replacements, for which CROW generates few variants. We identify three main challenges for diversification.

1) Constant computation We have observed that Souper searches for a constant replacement for more than 45% of the blocks of each program while constant values cannot be inferred. For instance, constant values cannot be inferred for memory load operations because CROW is oblivious to a memory model.

2) Combination computation The overlap between code blocks, mentioned in subsection 1.1.1, is a second factor that limits the number of unique variants. CROW can generate a high number of variants, but not all replacement combinations are necessarily unique. Let us illustrate the case with the example in Listing 1.2.

1.3.2 Properties for large diversification using CROW

We made a manual analysis of the programs that yield more than 100 unique variants to study the key properties of programs leveraging a high number of variants. This reveals one key reason that favors many unique variants: the programs include bounded loops. In these cases CROW synthesizes variants for the loops by unrolling them. Every time a loop is unrolled, the loop body is copied and moved as part of the outer scope of the loop. This creates a new, statically different, program. The number of programs grows exponentially with nested loops. On the other hand, CROW intensively tries to infer constants for all possible intermediate results in the input program, overcoming the limits of loop unrolls, this case is illustrated and discussed in a previous work of us **TODO** Superoptimization paper cite .

A second key factor for the synthesis of many variants relates to the presence of arithmetic. Souper, the synthesis engine used by CROW, is effective in replacing arithmetic instructions by equivalent instructions that lead to the same result. For example, CROW generates unique variants by replacing multiplications with additions or shift left instructions (Listing 1.5). Also, logical comparisons are replaced, inverting the operation and the operands (Listing 1.6). Besides, CROW is able to use overflow and underflow of integers to produce variants (Listing 1.7), using the intrinsics of underlying the computation model.

Listing 1.5: Diversification through arithmetic expression replacement.

Listing 1.6: Diversification through inversion of comparison operations.

Listing 1.7: Diversification through overflow of integer operands.

```
local.get 0      local.get 0
i32.const 2      i32.const 1
i32.mul          i32.shl
```

```
local.get 0    i32.const 11
i32.const 10   local.get 0
i32.gt_s       i32.le_s
```

```
i32.const 2    i32.const 2
i32.mul        i32.mul
               i32.const -2147483647
               i32.mul
```

1.3.3 Variant properties

Regarding the potential size overhead of the generated variants, we have compared the WebAssembly binary size of the diversified programs with their variants. The ratio of size change between the original program and the variants ranges from 82% (variants are smaller) to 125% (variants are larger) for **CROW prime** and **MEWE prime**. This limited impact on the binary size of the variants is good news because they are meant to save bandwidth when they become assets to distribute over network.

1.4 Conclusions

The proposed methodology is able to generate program variants that are syntactically different to their original versions. We have shown that CROW generates diversity among the binary code variants. We identified the properties that original programs should have to provide a handful number of variants using CROW. Besides, we enumerated the challenges faced to provide automatic diversification by retargeting a superoptimizer.

In the next chapter we evaluate the assessment of the generated variants answering to what extent the artificially programs are different from the original in terms of static difference, execution behavior and preservation.