



Software Diversification for WebAssembly

JAVIER CABRERA-ARTEAGA

Doctoral Thesis in Computer Science
Supervised by
Benoit Baudry and Martin Monperrus

Stockholm, Sweden, March 2024

KTH Royal Institute of Technology
School of Electrical Engineering and Computer Science
Division of Software and Computer Systems
TRITA-EECS-AVL-2020:4 SE-10044 Stockholm
ISBN 100- Sweden

Akademisk avhandling som med tillstånd av Kungl Tekniska högskolan framlägges
till offentlig granskning för avläggande av Teknologie doktorexamen i elektroteknik
i .

© Javier Cabrera-Arteaga , March 7th 2024

Tryck: Universitetsservice US AB

Abstract

WebAssembly, now the fourth officially recognized web language, enables web browsers to port native applications for the Web. Furthermore, WebAssembly has evolved into an essential element for backend scenarios such as cloud computing and edge computing. Therefore, WebAssembly finds use in a plethora of applications, including but not limited to, web browsers, blockchain, and cloud computing. Despite the emphasis on security since its design and specification, WebAssembly remains susceptible to various forms of attacks, including memory corruption and side-channels. Furthermore, WebAssembly has been manipulated to disseminate malware, particularly in cases of browser cryptojacking.

Web page resources, including those containing WebAssembly binaries, are predominantly served from centralized data centers in the modern digital landscape. Thousands of edge devices, in conjunction with browser clients, operate millions of identical WebAssembly instantiations every second. This phenomenon creates a highly predictable ecosystem, wherein potential attackers can anticipate behavior either in browsers or backend nodes. Such predictability escalates the potential impact of vulnerabilities within these ecosystems, paving the way for high-impact side-channel and memory attacks. For instance, a flaw in a web browser, instigated by a defective WebAssembly program, holds the potential to affect millions of users.

This thesis aims to bolster the security within the WebAssembly ecosystem through the introduction of Software Diversification methods and tools. Software Diversification is a strategy designed to augment the costs of exploiting vulnerabilities by making software unpredictable. The unpredictability within ecosystems can be diminished by automatically generating various program variants. These variants strengthen observable properties that are typically used to launch attacks, and in many instances, can completely eliminate such vulnerabilities.

This work introduces three tools: CROW, MEWE, and WASM-MUTATE. Each tool has been specifically designed to tackle a unique facet of Software Diversification. We present empirical evidence demonstrating the potential application of our Software Diversification methods to WebAssembly programs in two distinct ways: Offensive and Defensive Software Diversification. Our research into Offensive Software Diversification in WebAssembly unveils potential paths for enhancing the detection of WebAssembly malware. On the other hand, our experiments in Defensive Software Diversification show that WebAssembly programs can be hardened against side-channel attacks, specifically the Spectre attack.

Keywords: WebAssembly, Software Diversification, Side-Channels

Sammanfattning

LIST OF PAPERS

1. ***WebAssembly Diversification for Malware Evasion***
Javier Cabrera-Arteaga, Tim Toady, Martin Monperrus, Benoit Baudry
Computers & Security, Volume 131, 2023, 17 pages
<https://www.sciencedirect.com/science/article/pii/S0167404823002067>
2. ***Wasm-mutate: Fast and Effective Binary Diversification for WebAssembly***
Javier Cabrera-Arteaga, Nicholas Fitzgerald, Martin Monperrus, Benoit Baudry
Submitted to Computers & Security, under revision, 20 pages
<https://arxiv.org/pdf/2309.07638.pdf>
3. ***Multi-Variant Execution at the Edge***
Javier Cabrera-Arteaga, Pierre Laperdrix, Martin Monperrus, Benoit Baudry
Moving Target Defense (MTD 2022), 12 pages
<https://dl.acm.org/doi/abs/10.1145/3560828.3564007>
4. ***CROW: Code Diversification for WebAssembly***
Javier Cabrera-Arteaga, Orestis Floros, Oscar Vera-Pérez, Benoit Baudry, Martin Monperrus
Measurements, Attacks, and Defenses for the Web (MADWeb 2021), 12 pages
<https://doi.org/10.14722/madweb.2021.23004>
5. ***Superoptimization of WebAssembly Bytecode***
Javier Cabrera-Arteaga, Shrinish Donde, Jian Gu, Orestis Floros, Lucas Satabin, Benoit Baudry, Martin Monperrus
Conference Companion of the 4th International Conference on Art, Science, and Engineering of Programming (Programming 2021), MoreVMs, 4 pages
<https://doi.org/10.1145/3397537.3397567>
6. ***Scalable Comparison of JavaScript V8 Bytecode Traces***
Javier Cabrera-Arteaga, Martin Monperrus, Benoit Baudry
11th ACM SIGPLAN International Workshop on Virtual Machines and Intermediate Languages (SPLASH 2019), 10 pages
<https://doi.org/10.1145/3358504.3361228>

ACKNOWLEDGEMENT

Contents

List of Papers	iii
Acknowledgement	iv
Contents	1
 I Thesis	 4
1 Introduction	5
1.1 Predictability in WebAssembly ecosystems	7
1.2 Problems statements	9
1.3 Software Diversification	9
1.4 Summary of research papers.	10
 2 Background and state of the art	 13
2.1 WebAssembly	13
2.1.1 From source code to WebAssembly	14
2.1.2 WebAssembly’s binary format	18
2.1.3 WebAssembly’s runtime	19
2.1.4 WebAssembly’s control-flow	21
2.1.5 Security and reliability for WebAssembly	22
2.1.6 Open challenges	23
2.2 Software diversification	24
2.2.1 Automatic generation of software variants	25
2.2.2 Equivalence Checking	27
2.2.3 Variants deployment.	28
2.2.4 Measuring Software Diversification	29
2.2.5 Offensive or Defensive assessment of diversification	31

2.3	Open challenges for Software Diversification	31
3	Automatic Software Diversification for WebAssembly	33
3.1	CROW: Code Randomization of WebAssembly	34
3.1.1	Enumerative synthesis	35
3.1.2	Constant inferring	36
3.1.3	Exemplifying CROW	37
3.2	MEWE: Multi-variant Execution for WebAssembly	39
3.2.1	Multivariant call graph.	40
3.2.2	Exemplifying a Multivariant binary	40
3.3	WASM-MUTATE: Fast and Effective Binary Diversification for WebAssembly	43
3.3.1	WebAssembly Rewriting Rules	44
3.3.2	E-Graphs traversals	45
3.3.3	Exemplifying WASM-MUTATE	46
3.4	Comparing CROW, MEWE, and WASM-MUTATE	48
3.4.1	Security applications	51
4	Assessing Software Diversification for WebAssembly	53
4.1	Offensive Diversification: Malware evasion.	53
4.1.1	Cryptojacking defense evasion	54
4.1.2	Methodology	55
4.1.3	Results	57
4.2	Defensive Diversification: Speculative Side-channel protection . .	60
4.2.1	Threat model: speculative side-channel attacks	61
4.2.2	Methodology	62
4.2.3	Results	64
5	Conclusions and Future Work	69
5.1	Summary of technical contributions	69
5.2	Key results of the thesis	70
5.3	Future Work	71
5.3.1	Data augmentation for Machine Learning on WebAssembly programs.	71
5.3.2	Improving WebAssembly malware detection via canonicalization	72
5.3.3	Oneshot Diversification	73

<i>CONTENTS</i>	3
References	74
II Included papers	89
WebAssembly Diversification for Malware Evasion	91
Wasm-mutate: Fast and Effective Binary Diversification for WebAssembly	92
CROW: Code Diversification for WebAssembly	93
Multi-Variant Execution at the Edge	94
Superoptimization of WebAssembly Bytecode	95
Scalable Comparison of JavaScript V8 Bytecode Traces	96

Part I

Thesis